

[ZBD] Week 3 - PostgreSQL queries

Zacharczuk Jakub

jz418488

1 Zapytanie o pary ze wspólnym podmiotem [NEO4J]

1.1 Zapytanie z treści zadania

```
MATCH (a:Officer)-->(e:Entity)<--(b:Officer)
WITH a, b, COUNT(*) AS cnt
ORDER BY cnt DESC
RETURN a, b, cnt
LIMIT 10
```

Powyższe zapytanie działa bardzo długo (kilkadziesiąt sekund), po czym generuje błąd: "Neo.DatabaseError.General.UnknownError" z informacją "Java heap space". Z tego powodu chcemy ograniczyć powyższe zapytanie.

1.2 Nowe zapytanie

```
MATCH (a:Officer)-->(e:Entity)<--(b:Officer)
WHERE a.name STARTS WITH "A"
WITH a, b, COUNT(*) AS cnt
ORDER BY cnt DESC
RETURN a, b, cnt
LIMIT 10
```

1.3 Analiza wydajności nowego zapytania

Returned 10 rows in 830 ms.
Returned 10 rows in 485 ms.
Returned 10 rows in 428 ms.
Returned 10 rows in 426 ms.
Returned 10 rows in 422 ms.
Returned 10 rows in 446 ms.

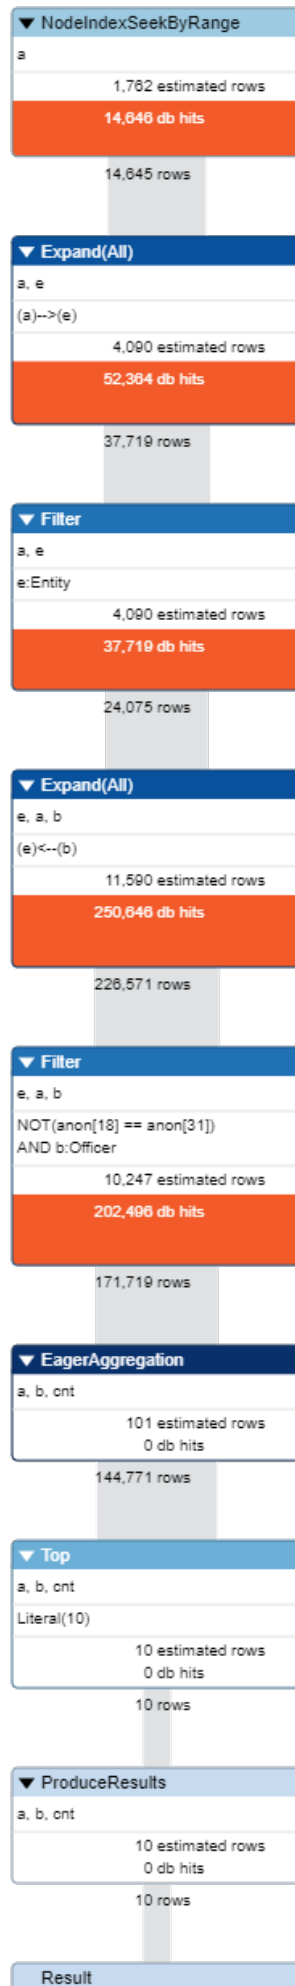
Wobec powyższych rezultatów można zauważyć, że pierwsze zapytanie wykonuje się znacznie dłużej niż kolejne. Najprawdopodobniej wynika to z faktu, że kolejne zapytania wykorzystują scachowany plan zapytania oraz pewne wiersze. Czasy po "rozgrzaniu maszyny" oscylują wokół 420ms.

1.4 Plan zapytania

Po wywołaniu poniższego zapytania:

```
PROFILE
MATCH (a:Officer)-->(e:Entity)<--(b:Officer)
WHERE a.name STARTS WITH "A"
WITH a, b, COUNT(*) AS cnt
ORDER BY cnt DESC
RETURN a, b, cnt
LIMIT 10
```

Otrzymujemy:



2 Zapytanie o pary ze wspólnym podmiotem [PostgreSQL]

2.1 Założenia:

W rozwiązaniu zdecydowałem się dodać poniższe założenia

- nie interesują nas takie pary Officers,

gdzie $a \rightarrow e \leftarrow a$

- chcemy pozbyć się duplikatów tzn.

$a \rightarrow e \leftarrow b$ to to samo co $b \rightarrow e \leftarrow a$

- nie musimy się ograniczać do Officers a zaczynających się na 'A'

2.2 Zapytanie:

2.2.1 Wybieramy tylko krawędzie wskazujące na Entity

```
SELECT start_id, end_id, name FROM edge
INNER JOIN entity ON entity.node_id = end_id
LIMIT 10;
```

2.2.2 Zawężamy do wierszy gdzie a jest typu Officer

```
SELECT start_id, officer.name, end_id, entity.name FROM edge
INNER JOIN entity ON entity.node_id = end_id
INNER JOIN officer ON officer.node_id = start_id
LIMIT 10;
```

2.2.3 Tworzymy krotki typu $(a:\text{Officer}) \rightarrow (e:\text{Entity}) \leftarrow (b:\text{Officer})$

```
SELECT e1.start_id, of1.name, e1.end_id, entity.name, e2.start_id FROM edge e1
INNER JOIN entity ON entity.node_id = e1.end_id
INNER JOIN officer of1 ON of1.node_id = start_id
INNER JOIN edge e2 ON e1.end_id = e2.end_id
INNER JOIN officer of2 ON of2.node_id = e2.start_id
LIMIT 10;
```

2.2.4 Grupujemy krotki po parach a, b i zliczamy liczbę wspólnych e $(a:\text{Officer}) \rightarrow (e:\text{Entity}) \leftarrow (b:\text{Officer})$

```
SELECT e1.start_id AS a_id, e2.start_id AS b_id, COUNT(entity.node_id) AS common_entities
FROM edge e1
INNER JOIN entity ON entity.node_id = e1.end_id
INNER JOIN officer of1 ON of1.node_id = start_id
INNER JOIN edge e2 ON e1.end_id = e2.end_id
INNER JOIN officer of2 ON of2.node_id = e2.start_id
GROUP BY a_id, b_id
ORDER BY common_entities DESC
LIMIT 10;
```

2.2.5 Aby uniknąć krotek opisanych w założeniach łączymy tylko takie a i b, gdzie $a.id < b.id$

```
...
INNER JOIN officer of2 ON of2.node_id = e2.start_id AND e2.start_id > e1.start_id
...
```

2.3 Ostateczna wersja zapytania

```
EXPLAIN ANALYZE
SELECT e1.start_id AS a_id, e2.start_id AS b_id, COUNT(entity.node_id) AS common_entities
FROM edge e1
    INNER JOIN entity ON entity.node_id = e1.end_id
    INNER JOIN officer of1 ON of1.node_id = start_id
    INNER JOIN edge e2 ON e1.end_id = e2.end_id
    INNER JOIN officer of2 ON of2.node_id = e2.start_id AND e2.start_id > e1.start_id
GROUP BY a_id, b_id
ORDER BY common_entities DESC
LIMIT 10;
```

2.4 Wynik zapytania:

a_id	b_id	common_entities
12215501	12215506	817
12211989	12217554	540
12205089	12205090	516
12138721	12215506	473
12132169	12175775	437
12118873	12140444	432
12130995	12197455	416
12096558	12140444	396
12204607	12204608	329
12215503	12215506	305

2.5 Wydajność

Aby zmierzyć wydajność wpisałem timing ON. W wyniku otrzymałem poniższe czasy:

Time: 1185.487 ms (00:01.185)

Time: 1218.000 ms (00:01.218)

Time: 1158.426 ms (00:01.158)

Time: 1196.862 ms (00:01.197)

Time: 1166.028 ms (00:01.166)

Wobec powyższych danych nie można stwierdzić że wygenerowanie planu zapytania skróciło czas działania. Być może restart kontenera i restart połączenia z bazą nie wystarczył do wyczyszczenia cache'a.

2.6 Plan zapytania:

Plan zapytania wygenerowałem dodając prefix EXPLAIN ANALYZE.

```
QUERY PLAN
-----
Limit  (cost=196869.56..196869.58 rows=10 width=16) (actual time=1525.732..1533.679
rows=10 loops=1)
  -> Sort  (cost=196869.56..198642.11 rows=709022 width=16) (actual
time=1511.888..1519.833 rows=10 loops=1)
    Sort Key: (count(entity.node_id)) DESC
    Sort Method: top-N heapsort  Memory: 25kB
  -> Finalize GroupAggregate  (cost=95918.74..181547.85 rows=709022 width=16)
(actual time=639.999..1379.559 rows=1638334 loops=1)
    Group Key: e1.start_id, e2.start_id
    -> Gather Merge  (cost=95918.74..170026.24 rows=590852 width=16)
(actual time=639.989..1064.875 rows=1647477 loops=1)
        Workers Planned: 2
        Workers Launched: 2
        -> Partial GroupAggregate  (cost=94918.71..100827.23
rows=295426 width=16) (actual time=606.875..805.361 rows=549159
loops=3)
```

```

Group Key: e1.start_id, e2.start_id
-> Sort (cost=94918.71..95657.28 rows=295426 width=12)
(actual time=606.858..661.858 rows=661267 loops=3)
    Sort Key: e1.start_id, e2.start_id
    Sort Method: external merge  Disk: 15208kB
    Worker 0:  Sort Method: external merge  Disk:
    12456kB
    Worker 1:  Sort Method: external merge  Disk:
    15208kB
-> Parallel Hash Join (cost=38756.77..63025.16
rows=295426 width=12) (actual time=254.795..409.997
rows=661267 loops=3)
    Hash Cond: (e1.end_id = entity.node_id)
    Join Filter: (e2.start_id > e1.start_id)
    Rows Removed by Join Filter: 785207
    -> Parallel Hash Join
        (cost=5721.02..20752.09 rows=280876 width=8)
        (actual time=29.584..73.810 rows=151719
        loops=3)
            Hash Cond: (e1.start_id = of1.node_id)
            -> Parallel Seq Scan on edge e1
                (cost=0.00..14293.76 rows=280876
                width=8) (actual time=0.014..17.942
                rows=224701 loops=3)
            -> Parallel Hash
                (cost=4479.34..4479.34 rows=99334
                width=4) (actual time=18.150..18.150
                rows=79467 loops=3)
                Buckets: 262144  Batches: 1
                Memory Usage: 11424kB
            -> Parallel Seq Scan on officer
                of1 (cost=0.00..4479.34
                rows=99334 width=4) (actual
                time=0.010..7.072 rows=79467
                loops=3)
        -> Parallel Hash (cost=28152.80..28152.80
        rows=280876 width=12) (actual
        time=161.610..161.612 rows=103119 loops=3)
            Buckets: 131072  Batches: 16  Memory
            Usage: 2016kB
            -> Parallel Hash Join
                (cost=12384.42..28152.80 rows=280876
                width=12) (actual time=69.748..145.226
                rows=103119 loops=3)
                Hash Cond: (e2.start_id =
                of2.node_id)
                -> Parallel Hash Join
                    (cost=6663.41..21694.48
                    rows=280876 width=12) (actual
                    time=16.391..90.740 rows=174330
                    loops=3)
                    Hash Cond: (e2.end_id =
                    entity.node_id)
                    -> Parallel Seq Scan on
                    edge e2
                        (cost=0.00..14293.76
                        rows=280876 width=8) (actual
                        time=0.031..20.594
                        rows=224701 loops=3)
                    -> Parallel Hash
                        (cost=5550.73..5550.73

```

```

rows=89014 width=4) (actual
time=15.313..15.313
rows=71211 loops=3)
    Buckets: 262144
    Batches: 1 Memory
    Usage: 10464kB
    -> Parallel Index
    Only Scan using
    entity_pkey on entity
    (cost=0.42..5550.73
    rows=89014 width=4)
    (actual
    time=0.023..5.524
    rows=71211 loops=3)
        Heap Fetches: 0
-> Parallel Hash
(cost=4479.34..4479.34 rows=99334
width=4) (actual
time=25.412..25.413 rows=79467
loops=3)
    Buckets: 262144 Batches: 1
    Memory Usage: 11424kB

-> Parallel Seq Scan on
officer of2
(cost=0.00..4479.34
rows=99334 width=4) (actual
time=8.549..15.224
rows=79467 loops=3)

```

Planning Time: 0.837 ms

JIT:

Functions: 121

Options: Inlining false, Optimization false, Expressions true, Deforming true

Timing: Generation 6.159 ms, Inlining 0.000 ms, Optimization 1.556 ms, Emission
37.315 ms, Total 45.030 ms

Execution Time: 1537.354 ms

3 Najkrótsza ścieżka [NEO4J]

3.1 Zapytanie

```

MATCH (a:Officer {name: "Emma Watson"})
MATCH (b:Officer {name: "APPLETON INVEST CAPITAL LIMITED"})
MATCH p=shortestPath((a)-[*]-(b))
RETURN p
LIMIT 10

```

3.2 Wydajność:

Returned 2 rows in 75 ms.

Returned 2 rows in 20 ms.

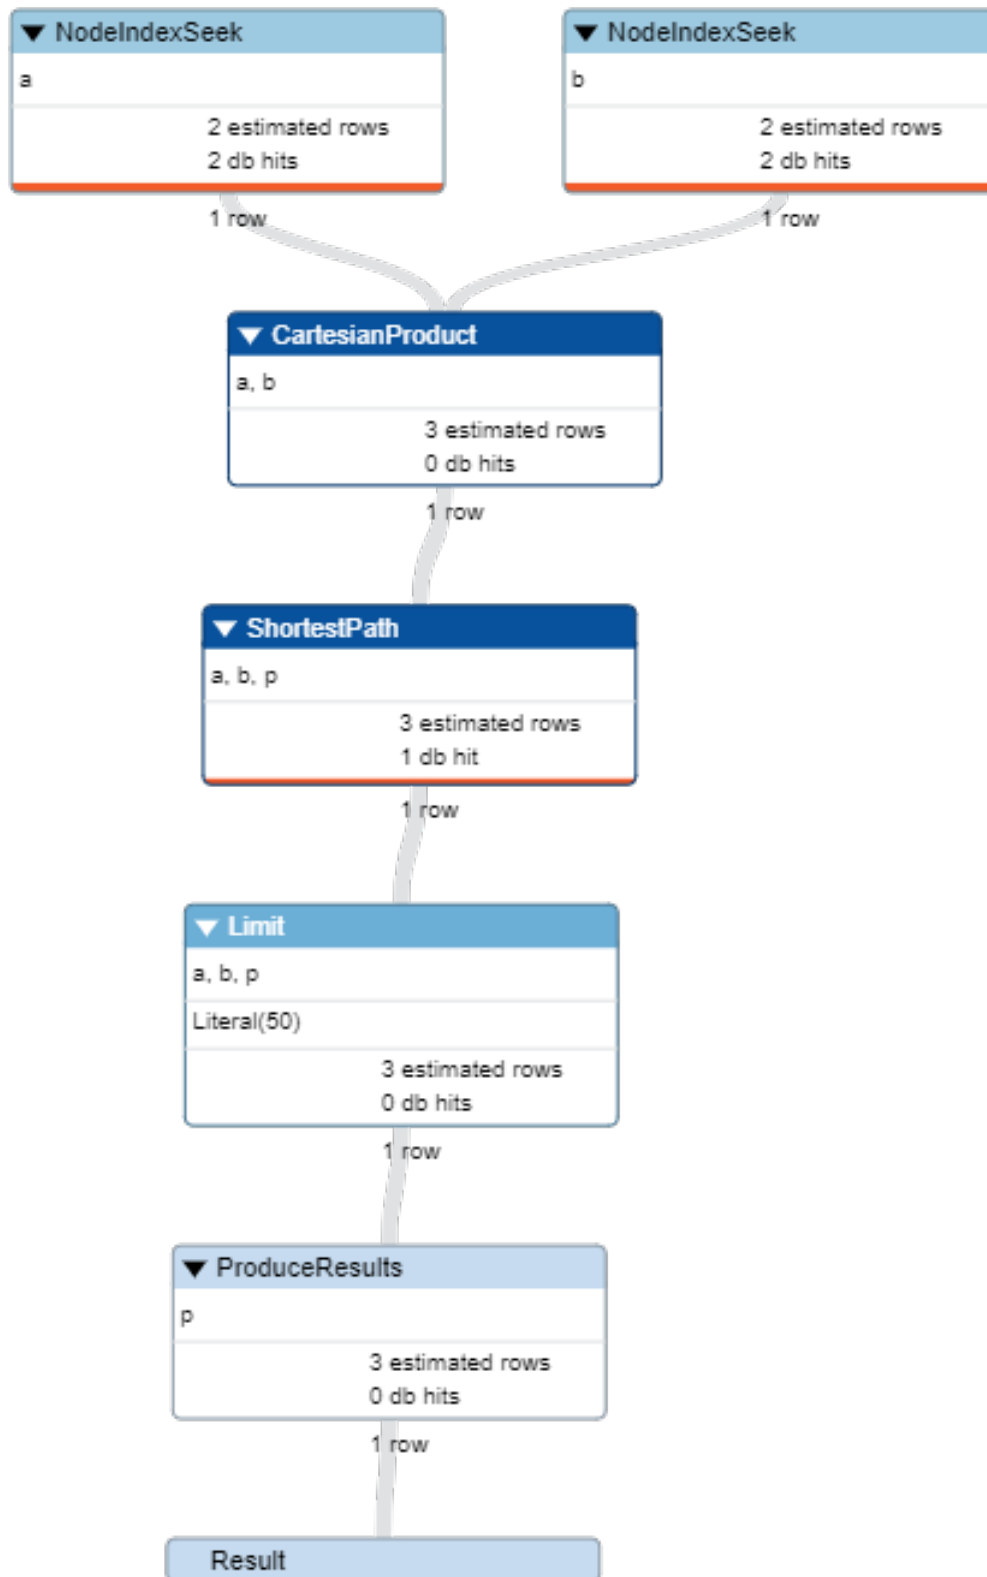
Returned 2 rows in 22 ms.

Returned 2 rows in 30 ms.

Returned 2 rows in 24 ms.

Ponownie widać, że kolejne wywołania tego samego zapytania działają znacznie szybciej.

3.3 Plan zapytania:



4 Najkrótsza ścieżka [PostgreSQL]

4.1 Zapytanie:

id: 12126782 - Emma Watson

id: 12132897 - APPLETON INVEST CAPITAL LIMITED

```
WITH RECURSIVE generate_path (id, depth, visited) AS (  
    SELECT start_id AS id,  
           0 AS depth,  
           ARRAY[start_id] AS visited  
    FROM edge  
    WHERE start_id = 12126782  
    UNION ALL (  
        (WITH generate_path(id, depth, visited) AS (TABLE generate_path)  
        SELECT e.end_id, depth + 1, visited || e.end_id  
        FROM generate_path gp, edge e  
        WHERE gp.id = e.start_id AND NOT e.end_id = ANY(visited)  
              AND depth < 7  
        UNION ALL  
        SELECT e.start_id, depth + 1, visited || e.start_id  
        FROM generate_path gp, edge e  
        WHERE gp.id = e.end_id AND NOT e.start_id = ANY(visited)  
              AND depth < 7  
        ))  
    )  
SELECT visited  
FROM generate_path  
WHERE id = 12132897  
LIMIT 1;
```

4.2 Wynik

```
visited  
-----  
{12126782,10152535,11011539,10155950,12132897}  
Time: 456.378 ms
```

4.3 Plan zapytania (wygenerowany jak wcześniej)

```
QUERY PLAN  
-----  
Limit (cost=431484.07..431488.56 rows=1 width=32) (actual time=663.135..663.193  
rows=1 loops=1)  
  CTE generate_path  
    -> Recursive Union (cost=1000.00..431484.07 rows=5380 width=40) (actual  
time=32.934..644.860 rows=5095 loops=1)  
      -> Gather (cost=1000.00..15996.95 rows=10 width=40) (actual  
time=32.932..33.090 rows=2 loops=1)  
        Workers Planned: 2  
        Workers Launched: 2  
        -> Parallel Seq Scan on edge (cost=0.00..14995.95 rows=4 width=40) (actual  
time=15.933..16.021 rows=1 loops=3)  
          Filter: (start_id = 12126782)  
          Rows Removed by Filter: 224700  
      -> Append (cost=4.66..41537.95 rows=537 width=40) (actual time=79.179..152.762  
rows=1273  
loops=4)  
        CTE generate_path  
          -> WorkTable Scan on generate_path generate_path_1 (cost=0.00..2.00
```



```

    rows=100 width=40) (actual time=0.000..0.049 rows=1272 loops=4)
-> Hash Join (cost=2.66..20766.99 rows=379 width=40) (actual
time=70.864..84.878 rows=1270 loops=4)
    Hash Cond: (e.start_id = gp.id)
    Join Filter: (e.end_id <> ALL (gp.visited))
    Rows Removed by Join Filter: 2
-> Seq Scan on edge e (cost=0.00..18226.02 rows=674102 width=8)
(actual time=0.002..41.239 rows=674102 loops=4)
-> Hash (cost=2.25..2.25 rows=33 width=40) (actual time=0.377..0.377
rows=1272 loops=4)
    Buckets: 8192 (originally 1024) Batches: 1 (originally 1)
    Memory Usage: 461kB
-> CTE Scan on generate_path gp (cost=0.00..2.25 rows=33
width=40) (actual time=0.003..0.246 rows=1272 loops=4)
    Filter: (depth < 7)
-> Hash Join (cost=2.66..20760.91 rows=158 width=40) (actual
time=53.464..67.817 rows=3 loops=4)
    Hash Cond: (e_1.end_id = gp_1.id)
    Join Filter: (e_1.start_id <> ALL (gp_1.visited))
    Rows Removed by Join Filter: 1
-> Seq Scan on edge e_1 (cost=0.00..18226.02 rows=674102 width=8)
(actual time=0.002..31.782 rows=505622 loops=4)
-> Hash (cost=2.25..2.25 rows=33 width=40) (actual time=0.279..0.279
rows=1272 loops=4)
    Buckets: 8192 (originally 1024) Batches: 1 (originally 1)
    Memory Usage: 461kB
-> CTE Scan on generate_path gp_1 (cost=0.00..2.25 rows=33
width=40) (actual time=0.002..0.149 rows=1272 loops=4)
    Filter: (depth < 7)
-> CTE Scan on generate_path (cost=0.00..121.05 rows=27 width=32) (actual
time=645.934..645.934 rows=1 loops=1)
    Filter: (id = 12132897)
    Rows Removed by Filter: 5094
Planning Time: 0.184 ms
JIT:
  Functions: 50
  Options: Inlining false, Optimization false, Expressions true, Deforming true
  Timing: Generation 3.228 ms, Inlining 0.000 ms, Optimization 1.124 ms, Emission 20.760 ms,
  Total 25.112 ms
Execution Time: 666.105 ms
(39 rows)

Time: 667.072 ms

```

4.4 Analiza

Najprawdopodobniej z powodu takiego jak w punkcie 2, kolejne wywołania nie zmieniają czasu działania zapytania. Dodatkowo w zapytaniu zastosowałem sztywny limit na $\text{depth} < 7$, gdyż zwiększenie tej wartości powoduje znaczne wydłużenie działania zapytania (tzn. ponad minutę).

5 Wnioski

Bazę danych należy wybrać pod kątem wykonywanych na niej operacji. W operacjach typu 1 znacznie lepiej spisał się Postgres, natomiast w zapytaniach grafowych znaczną przewagę ma dedykowany takim zapytaniom Neo4j.