

# lab03 bonus

Zacharczuk Jakub

October 2021

## 1 BFS do Panama Papers. Pierwsze podejście:

```
WITH RECURSIVE generate_path (queue_it, visited, queue) AS (  
  SELECT 1 AS queue_it,  
         ARRAY[0] AS visited,  
         ARRAY[start_id] AS queue  
  FROM edge  
  WHERE start_id = 12126782 AND end_id = 10152535  
  UNION ALL (  
    (WITH generate_path(queue_it, visited, queue) AS (TABLE generate_path)  
     SELECT queue_it + 1, visited || queue[queue_it], queue || (  
SELECT ARRAY(  
  SELECT e.nd  
  FROM generate_path gp, (SELECT start_id AS st, end_id AS nd  
  FROM edge  
  UNION ALL (  
SELECT end_id AS st, start_id AS nd  
FROM edge)) e  
  WHERE gp.queue[queue_it] = e.st AND NOT e.nd = ANY(visited)  
        AND NOT e.nd = ANY(queue)  
  )  
)  
  FROM generate_path gp  
  WHERE queue_it < 200  
  ))  
)  
SELECT queue_it, cardinality(visited), cardinality(queue)  
FROM generate_path  
LIMIT 200;
```

Time: 4845.826 ms (00:04.846)

W powyższym przykładzie bardzo nieefektywnie, pobieramy pierwszy element z kolejki, kopiując ją i po każdym przejściu powiększając ją.

## 2 Tu jest znacznie lepiej:

```
WITH RECURSIVE generate_path (queue_it, visited, queue) AS (  
  SELECT 1 AS queue_it,  
         ARRAY[0] AS visited,  
         ARRAY[start_id] AS queue  
  FROM edge  
  WHERE start_id = 12126782 AND end_id = 10152535  
  UNION ALL (  
    (WITH generate_path(queue_it, visited, queue) AS (TABLE generate_path)  
     SELECT queue_it + 1, visited || queue, (  
SELECT ARRAY(  
  SELECT e.nd  
  FROM generate_path gp, (SELECT start_id AS st, end_id AS nd  
  FROM edge  
  UNION ALL (  
SELECT end_id AS st, start_id AS nd  
FROM edge)) e  
  WHERE e.st = ANY(queue) AND NOT e.nd = ANY(visited)  
        AND NOT e.nd = ANY(queue)  
  )  
)  
  FROM generate_path gp  
  ))  
)  
SELECT queue_it, cardinality(visited), cardinality(queue)  
FROM generate_path  
LIMIT 6;
```

queue_it	cardinality	cardinality
1	1	1
2	2	2
3	4	3
4	7	2539
5	2546	2659
6	5205	8613

(6 rows)

Time: 22849.114 ms (00:22.849)

## 3 Uniknięcie konieczności łączenia tabeli krawędzi

W celu przyspieszenia utworzyłem nową tabelę w bazie danych:

```
CREATE TABLE undirected_edges as (  
  SELECT start_id, end_id  
  FROM edge  
  UNION ALL (  
    SELECT end_id, start_id  
    FROM edge  
  )  
);
```

```

WITH RECURSIVE generate_path (depth, visited, queue) AS (
    SELECT 0 AS depth,
           ARRAY[0] AS visited,
ARRAY[node_id] AS queue
    FROM officer
    WHERE node_id = 12126782
    UNION ALL (
        (WITH generate_path(depth, visited, queue) AS (TABLE generate_path)
        SELECT depth + 1, visited || queue, (
SELECT ARRAY(
    SELECT e.end_id
    FROM generate_path gp, undirected_edges e
    WHERE e.start_id = ANY(queue) AND NOT e.end_id = ANY(visited) AND NOT e.end_id = ANY(queue)
    GROUP BY e.end_id
    )
    )
        FROM generate_path gp
    ))
)
SELECT depth, cardinality(visited) AS visited_nodes, cardinality(queue) AS queue_size
FROM generate_path
LIMIT 6;

```

depth	visited_nodes	queue_size
0	1	1
1	2	2
2	4	3
3	7	2539
4	2546	2155
5	4701	7098

(6 rows)

Time: 20435.161 ms (00:20.435)

#### 4 To jest najlepsze co mi się udało osiągnąć: (znajduje się też w pliku query.txt)

```

WITH RECURSIVE generate_path (depth, visited, queue) AS (
    SELECT 0 AS depth,
           ARRAY[0] AS visited,
ARRAY[node_id] AS queue
    FROM officer
    WHERE node_id = 12126782
    UNION ALL (
        SELECT depth + 1, visited || queue, (
SELECT ARRAY(
    SELECT e.end_id
    FROM unnest(gp.queue) elem_id
    INNER JOIN undirected_edges e
    ON elem_id = e.start_id AND NOT e.end_id = ANY(visited)
    GROUP BY e.end_id
    )
    )
        FROM generate_path gp
    )
)
SELECT depth, cardinality(visited) AS visited_nodes, cardinality(queue) AS queue_size
FROM generate_path
LIMIT 6;

```

depth	visited_nodes	queue_size
0	1	1
1	2	2
2	4	3
3	7	2539
4	2546	2155
5	4701	7098

(6 rows)

Time: 13643.054 ms (00:13.643)

## 5 Plan powyższego zapytania znajduje się w pliku plan.txt

## 6 Wnioski

Najwięcej czasu zajmuje przeglądanie czy element jest już w kolejce oraz czy był już odwiedzony. Nie wymyśliłem sprytnego sposobu jak można byłoby powyższe zapytanie przyspieszyć.